

Cyber Underwriting Workbench — User Research

Birlasoft · Research lock 2026-07-20 · Secondary only (no primary interviews yet)

1. Observations

Cyber underwriting still runs on self-attested applications. Applicants describe policies; underwriters know the live environment often differs (for example “endpoint detection and response everywhere” versus deployment gaps).

Specialty desks burn time hunting information — often half the day searching, not deciding. Soft market: more competition, rate pressure, brokers’ market. Desks must move faster and sometimes concede on price, terms, or controls.

Portfolio risk sits behind each quote: shared cloud, identity, software-as-a-service, and managed service provider spines. One vendor incident can hit many policies; concentration at quote time is rarely clear.

Caveat: synthesis from open industry sources — not Birlasoft managing general agent primary interviews.

2. Insights

The job is **reconcile claim versus signal, then decide and defend** under time and price pressure — not complete another questionnaire.

Intake automation alone will not differentiate. Without attest-versus-signal, the proof of concept is “underwriting workbench with cyber labels.” Soft market rewards tools that keep discipline *and* speed: appetite, refer / decline rationale, audit trail.

Do not rebuild BitSight / CyberCube or become Coalition / At-Bay. Birlasoft fit = **intellectual-property accelerator**: configurable Decision Workbench clients personalize (rules, connectors, branding, policy administration system).

3. Recommendations

Fund Band A proof of concept: messy intake → gap board → explainable Quote / Refer / Decline → dossier / audit → policy administration stub (+ pricing tiers + accumulation alert). Partner or defer continuous telemetry, catastrophe engines, rate adequacy.

Measure: time to decision-ready dossier; percent of submissions with attestation conflicts caught before quote; percent of decisions with auditable rationale; live appetite-rule swap for personalization demo.

Before scaling: eight to twelve primary cyber underwriter interviews at target managing general agents. Later roadmap: continuous posture, deeper accumulation, claims / incident-response learning.

4. Competitor outcomes (summary)

Detail (two pages or less): `COMPETITOR-RESEARCH.md` .

	Takeaway
a. Market signals	Pay for control-documented speed; signal / catastrophe / managing general agent stack is settled — orchestrate, do not rebuild; book-watch language matters; agentic triage is baseline narrative
b. Agentic architecture	Look like Horizon 2: specialist stages + orchestrator + human-in-the-loop cite-back. Not Horizon 3 autonomy in the first ship
c. Feature benchmark	Win: gap + appetite + dossier + accumulation stub + policy administration. Partner: ratings / catastrophe. Out: insurance-plus-security stack. Parity: intake / completeness / shared record

Detail: `COMPETITOR-RESEARCH.md` · `SECONDARY-RESEARCH-PAIN-POINTS.md` · `FEATURE-LISTS-MVP-MOON.md`